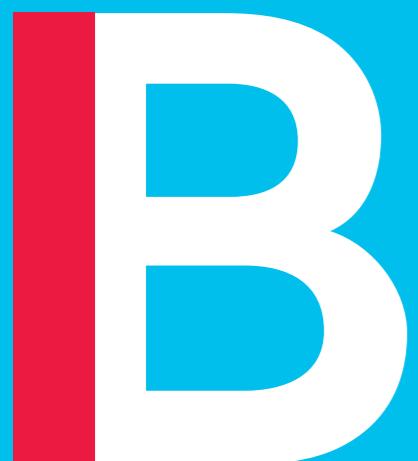


Machine Learning in Finance

Lecture 5
Practical Implementation : Word Vectors



Arnaud de Servigny & Hachem Madmoun

Outline:

- Introducing the Problem
- Word Embedding Methods
 - The GloVe approach
 - The Word2vec approach
- Programming Session: Implementation of the GloVe approach

Part 1 : Introducing the problem

Why do we need vectors to represent words ?

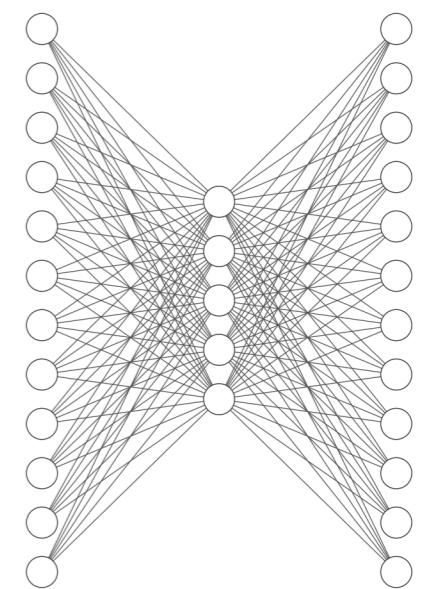
- We are dealing with data in the form of a corpus of sentences and want to perform a classification task for instance.
- We obviously can't feed words to a model. A model can only handle numbers.
- The question is : How do we represent the words of our corpus in a way that can be feeded in a Machine Learning Algorithm ?
- It's clearly an Unsupervised Learning task.

DATA

- Document 1 : « The sole evidence it is possible to produce that anything is desirable is that people actually desire it. »
- Document 2 : « In law a man is guilty when he violates the rights of others. In ethics he is guilty if he only thinks of doing so. »
- Document 3 : « Always recognize that human individuals are ends, and do not use them as means to your end. »
- ...
- Document N : « Justice is a name for certain moral requirements, which, regarded collectively, stand higher in the scale of social utility and are therefore of more paramount obligation than any others. »



Model



Review : Words as discrete symbols:

- What we have seen so far (in Lecture 5) is the possibility to turn each word into a discrete symbol.
- For that, we create a dictionary to map each word present in our corpus to a unique discrete index.

DATA

- Document 1 : « The sole evidence it is possible to produce that anything is desirable is that people actually desire it. »
- Document 2 : « In law a man is guilty when he violates the rights of others. In ethics he is guilty if he only thinks of doing so. »
- Document 3 : « Always recognize that human individuals are ends, and do not use them as means to your end. »
- ...
- Document N : « Justice is a name for certain moral requirements, which, regarded collectively, stand higher in the scale of social utility and are therefore of more paramount obligation than any others. »

Code

```
index = 1
word_index = {}
for sentence in sentences:
    for word in sentence:
        if word not in word_index:
            word_index[word] = index
        index += 1
```

word_index =

{

 'the' : 1 ,

 'sole' : 2 ,

 'evidence' : 3 ,

 ... : ...

 'any' : 934233

}

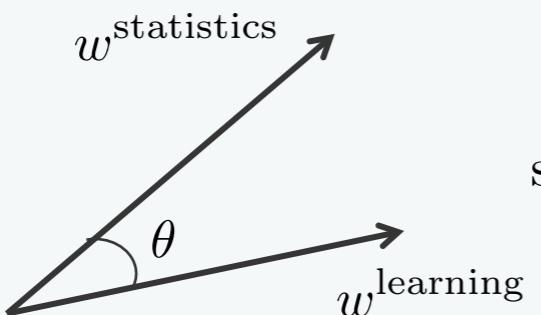
From discrete symbols to one hot vectors:

- | Corpus | Discretize via word_index |
|---------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none">• Document 1• Document 2• ⋮• Document n• ⋮• Document N | <ul style="list-style-type: none">• Document 1 : [23, 43, 12, ..., 2343, 1]• Document 2 : [12, 1, 23453, ..., 123]• ⋮• Document n : [1234, 1, 23]• ⋮• Document N : [1, 1232, ..., 12322] |
- After the first pre-processing step, we end up with the following lists of integers representing the words:
 - Instead of representing a word by its index in the **word_index** dictionary. It is strictly equivalent to represent it as a vector of size V (where V is the size of vocabulary, i.e the number of distinct words in the whole corpus) with 1 in the index position and zeros in all the other positions.
 - Example : Let's suppose the word « equity » is of index 134 and V = 100 000.
 - Then, the word « equity » will be represented by the following vector of size V:
$$[0, \dots, 0, 1, 0, \dots, 0]$$


position 134
 - We call this vector a **one hot vector**.

Limitations of one hot vectors:

- The **One-hot-vector** is the easiest way to represent words as vectors.
- In this type of encoding, each word is as a completely independent entity and there isn't any notion of similarity between words, even if they have the same meaning.
- One way of measuring the similarity between two vectors is to use the dot product.
- The dot product is just the cosine similarity:



$$\text{similarity}(w^{\text{statistics}}, w^{\text{learning}}) = \cos(\theta) = \frac{\langle w^{\text{statistics}}, w^{\text{learning}} \rangle}{\|w^{\text{statistics}}\| \|w^{\text{learning}}\|}$$

- In the case, the two words “statistics” and “learning” will have a similarity of zero, even though they are related to each other.

$$\langle w^{\text{statistics}}, w^{\text{learning}} \rangle = (w^{\text{statistics}})^T w^{\text{learning}} = [0 \ 1 \ 0 \ \dots \ 0] \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix} = 0$$

Part 2 : Word Embedding Methods

Creating Word Embedding

- We need to find a subspace that encodes the relationships between words.
- As there are millions of tokens in any language, we can try to reduce the size of this space from \mathbb{R}^V (where V is the vocabulary size) to some D-dimensional space (such that $D \ll V$) that is sufficient to encode all semantics of the language.
- Each dimension would encode some meaning (such as tense, count, gender ...)
- We are going to introduce 2 approaches:
 - **GloVe** (Global Vectors for Word Representation, Pennington, Socher and Manning, 2014).
 - The **Word2vec approach**: introduced by Mikolov, Sutskever, et al. (2013).
- Both algorithms take their inspiration from an English linguist, named John Ruper Firth, known for his famous quotation:

« *You shall know a word by the company it keeps* » (J.R. Firth 1957:11)

The GloVe approach – Introduction –

- **GloVe** (Global Vectors for Word Representation) is an unsupervised algorithm, developed at the Stanford NLP lab, that learns embedding vectors from word-word co-occurrence statistics.
- The GloVe algorithm consists in applying **Matrix factorization methods** on a matrix summarizing the **co-occurrence statistics** of the corpus.
- The entry X_{ij} of the matrix of co-occurrence counts X represents the number of times the word j occurs in the context of word i , which suggest the definition of a context size (or window size).
- Example with i = index of the word « enjoyed ». We append the index of word « I » twice.

$\overbrace{\quad \quad \quad \quad \quad}^{window\ size=1}$
 I enjoyed the research project.
 $X_{ij} += 1$

$\overbrace{\quad \quad \quad \quad \quad}^{window\ size=1}$
 I enjoyed NLP
 $X_{ij} += 1$

The GloVe approach – The co-occurrence matrix –

- Let us create the co-occurrence matrix on a simple corpus composed of 3 documents and a vocabulary size of 10 tokens. $V = 10$. So, the co-occurrence matrix is of shape (V, V)
- The corpus:
 - Document 1: I enjoyed the research project .
 - Document 2: : I like Deep Learning .
 - Document 3: I enjoyed NLP .
- The final co-occurrence matrix:

$$X = \begin{pmatrix} & & & & & & & & & & & V \\ & & & & & & & & & & & \\ I & 0 & 2 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ enjoyed & 2 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ the & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ research & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ project & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ like & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ Deep & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ Learning & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ NLP & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ . & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 1 & 0 & 0 \end{pmatrix}$$

V

V


The GloVe approach – SVD based methods –

- To create **embedding vectors** from the **co-occurrence matrix**, one approach can be to use a **Singular Value decomposition (SVD)** of the co-occurrence matrix:

$$X = W_1 \Omega W_2^T$$

$(V \times V) \quad (V \times V) \quad (V \times V) \quad (V \times V)$

- Then, we reduce the dimensionality by selecting the first D singular vectors (with $D \ll V$)

$$\hat{X} = \hat{W}_1 \hat{\Omega} \hat{W}_2^T$$


$(V \times V) \quad (V \times D) \quad (D \times D) \quad (D \times V)$

- Let $\Omega = \text{diag}(\omega_1, \dots, \omega_V)$, such that $\omega_1 > \omega_2 > \dots > \omega_V$

- We select D so that we can capture the desired amount of variance we want:

$$\frac{\sum_{i=1}^D \omega_i}{\sum_{i=1}^V \omega_i}$$

The GloVe approach – Matrix Factorization instead of SVD –

- The SVD approach does not work well in practice for several reasons:
 - The dimensions of the matrix change very often (new words are added very frequently and the corpus changes in size).
 - The matrix is extremely sparse (i.e, it contains a lot of zero values) since most words do not co-occur.
 - The matrix is very high dimensional as the vocabulary size is usually huge.
- We are going to introduce another way of performing the factorization: **Matrix Factorization Methods** are widely used for generating **meaningful and low-dimensional word representation**.
 - In the GloVe approach, since non-zero values are very large, we factorize the logarithm of X (denoted $\log X$) instead of factorizing X .
 - Remark: Obviously, as we can't apply the logarithm function on the entries with a zero value, we add 1 to all the element of the matrix before applying the logarithm).

$$\forall (i, j) \in V^2 \quad X_{ij} \leftarrow X_{ij} + 1$$

- We want to factorize $\log X$ into 2 matrices: $\log X \approx W\tilde{W}^T$
- We want to estimate $W, \tilde{W} \in \mathbb{R}^{V \times D}$ with $D \ll V$

Matrix Factorization for Collaborative Filtering

- Let us introduce the concept of Matrix Factorization in the context of **Collaborative Filtering**.
- Let us imagine that we have users who rate movies on some platform.
 - The number of users is N
 - The number of movies is K
 - Each rating is a real number

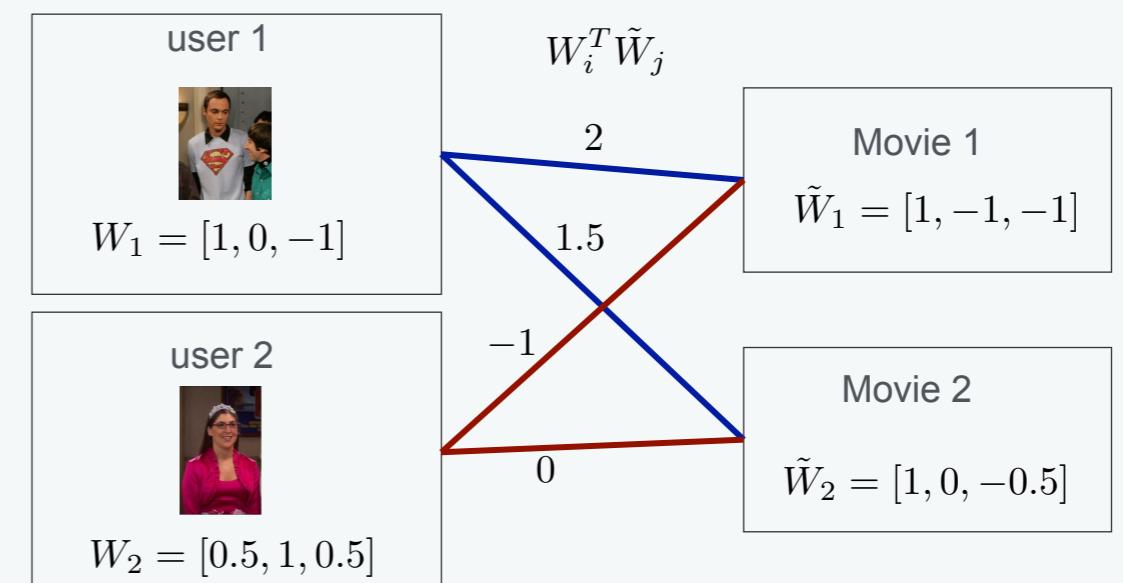
	Movie 1	...	Movie k	...	Movie K
User 1	-2				-1
:					
User n	5		4		
:					
User N	1				

- We end up with a very sparse matrix R of shape (N, K) (since most users have seen very few movies).
- Collaborative Filtering is the concept of using other people's rating to make a rating prediction about a movie for a user who has never seen it.

Matrix Factorization for Collaborative Filtering

- In order to approximate R , we factorize the matrix into two matrices $R \approx \hat{R} = W\tilde{W}^T$
 - We want to estimate the parameters W and \tilde{W} by minimizing $J = \sum_{i=1}^N \sum_{j=1}^K (R_{ij} - \hat{R}_{ij})^2 = \sum_{i=1}^N \sum_{j=1}^K (R_{ij} - W_i^T \tilde{W}_j)^2$
- $$W = \begin{pmatrix} - & W_1 & - \\ \vdots & \vdots & \vdots \\ - & W_N & - \end{pmatrix} \in \mathbb{R}^{N \times D}$$
- $$\tilde{W} = \begin{pmatrix} - & \tilde{W}_1 & - \\ \vdots & \vdots & \vdots \\ - & \tilde{W}_K & - \end{pmatrix} \in \mathbb{R}^{K \times D}$$
- Each row i of the W matrix is a D -dimensional vector representing the user i . Each dimension encodes a latent meaningful information about the user.
 - Each row j of the \tilde{W} matrix is a D -dimensional vector representing the movie j . Similarly, each dimension encodes a meaningful information about the movie.

- As an example: let us consider $D=3$ latent dimensions:
 - Sci-fi
 - Comedy
 - Romance



The GloVe approach

- In the GloVe approach, we are going to approximate the logarithm of the co-occurrence matrix by using the same factorization method.
- We also add a bias term for the matrix W and a bias term for \tilde{W}


$$\forall (i, j) \in \{1, \dots, V\}^2 \quad \log X_{ij} \approx W_i^T \tilde{W}_j + b_i + \tilde{b}_j$$

- The parameters of the model:

$$W = \begin{pmatrix} - & W_1 & - \\ \vdots & \vdots & \vdots \\ - & W_V & - \end{pmatrix} \in \mathbb{R}^{V \times D} \quad \tilde{W} = \begin{pmatrix} - & \tilde{W}_1 & - \\ \vdots & \vdots & \vdots \\ - & \tilde{W}_V & - \end{pmatrix} \in \mathbb{R}^{V \times D} \quad b = \begin{pmatrix} b_1 \\ \vdots \\ b_V \end{pmatrix} \in \mathbb{R}^V \quad \tilde{b} = \begin{pmatrix} \tilde{b}_1 \\ \vdots \\ \tilde{b}_V \end{pmatrix} \in \mathbb{R}^V$$

- The cost function of the weighted least squares regression model is:

$$J = \sum_{i=1}^V \sum_{j=1}^V f(X_{ij})(\log X_{ij} - W_i^T \tilde{W}_j - b_i - \tilde{b}_j)^2$$



- The weights $f(X_{ij})$ are added because we consider that rare occurrences are noisy and carry less information than the more frequent ones.

The Word2vec approach – The idea –

- To create word embedding using the Word2vec approach, the idea is to define a word by the context of this word in all the documents of the training corpus.
- For instance, let us consider the word « economy ». Obviously, the word « economy » is not going to appear in the same context as the word « rock ».

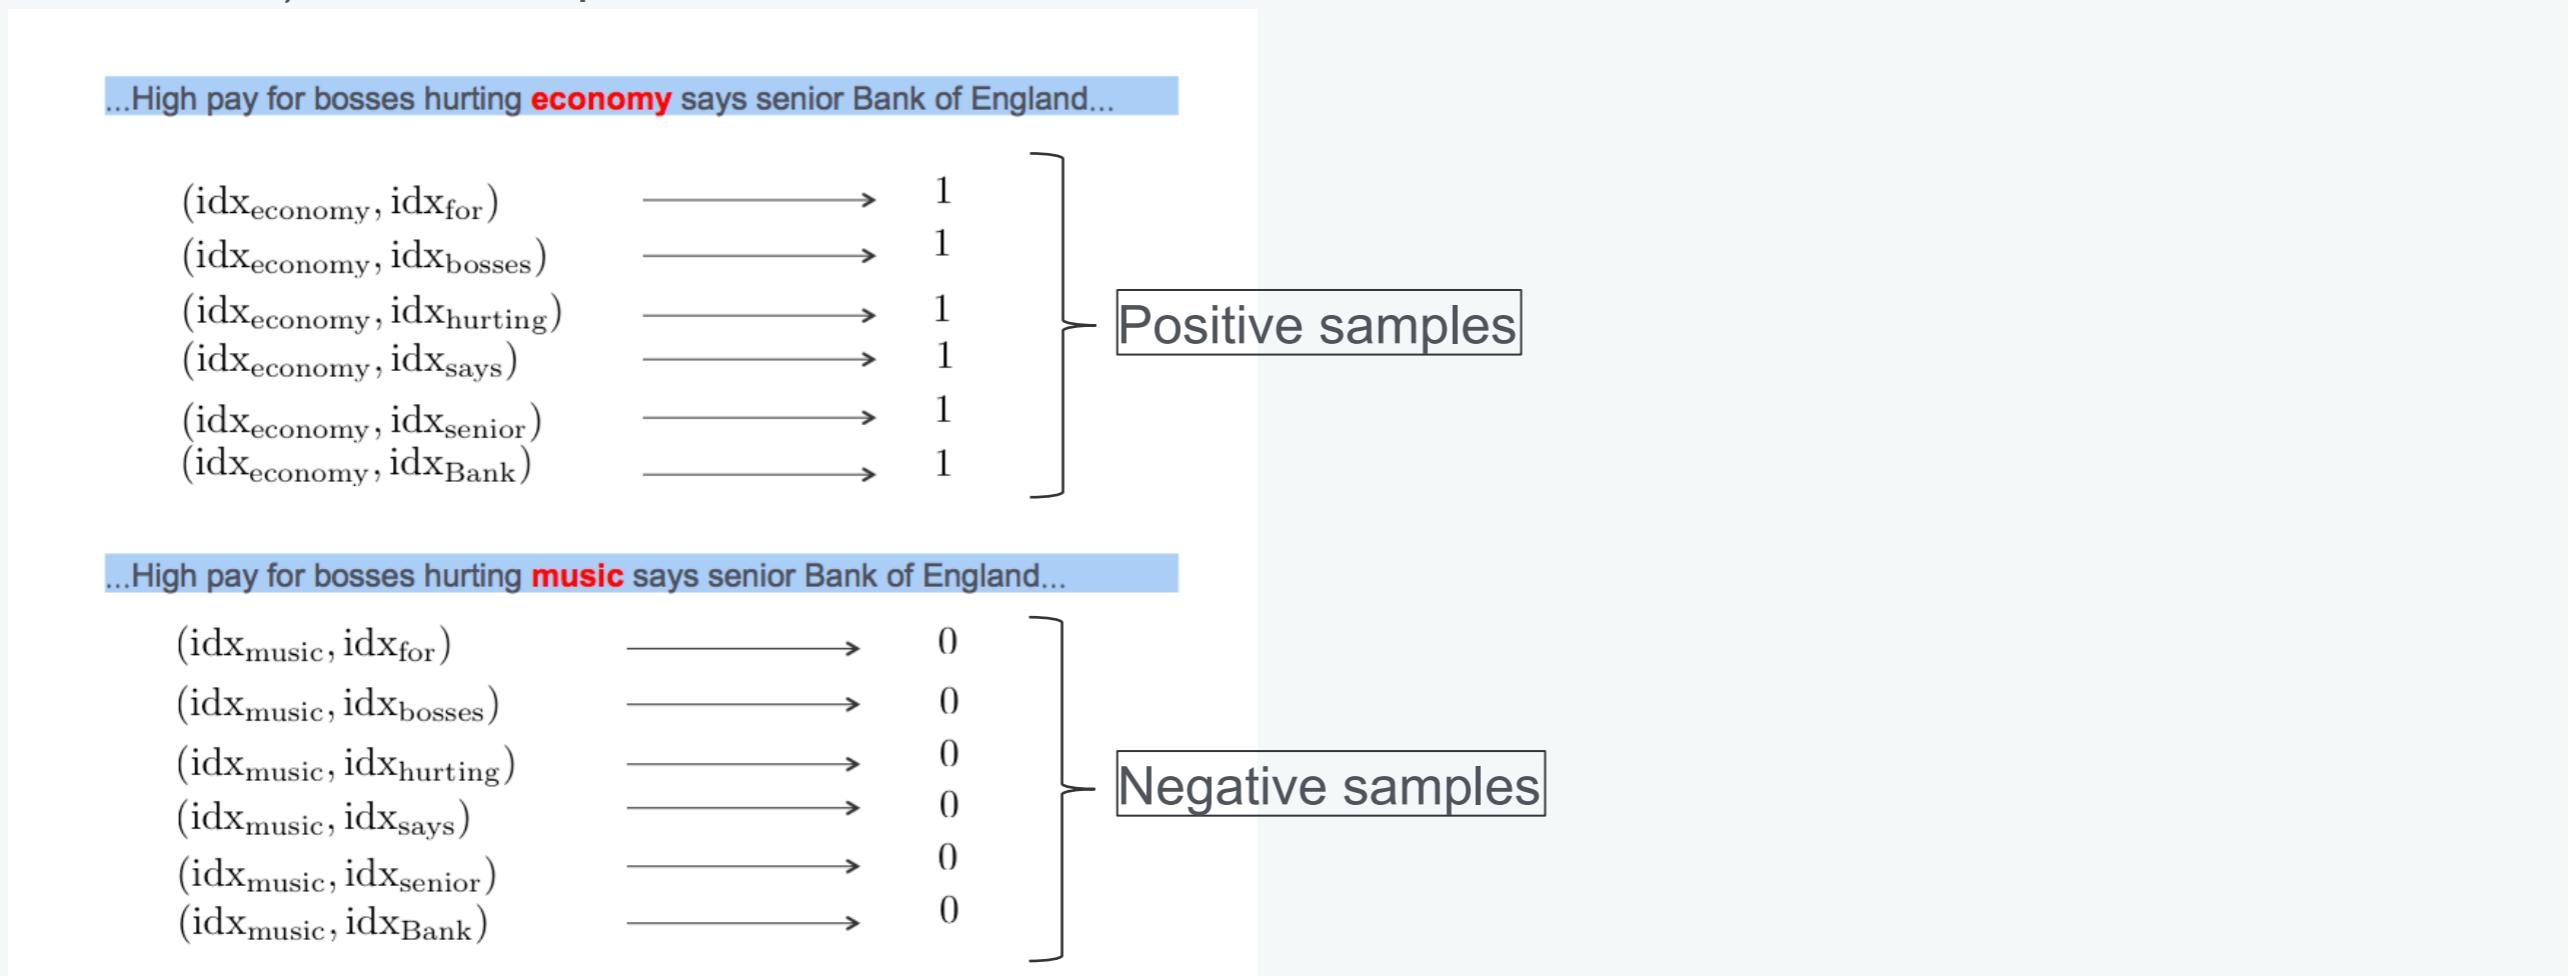
...dire consequences for the UK **economy**, even as markets were rocked...
...High pay for bosses hurting **economy** says senior Bank of England...
...Mervyn King believes the world **economy** will soon face another crash...



- The Word2vec approach consists in creating word embedding vectors by using a shallow neural network in order to:
 - Predict the center word (« economy » in our example ») from the context words. It's called The Continuous Bag of Words method (CBOW)
 - Predict the context words from the center word. It's called the Skipgram method.
- In our implementation we are going to focus on the Skipgram method with negative sampling.

The Word2vec approach – The data –

- Let us consider a window size of D.
- For each center word in our corpus, we have a list of $2*D$ context words associated with this center word (the ones on the right and the ones on the left).
- We can then define $2*D$ couples of (center word, context word) as shown in the figure with « economy » as a center word. These couples are associated with a label 1.
- By sampling a random word in the corpus, we can create other false couples of (center word, context word). These couples are associated with a label 0.



The Word2vec approach – The Forward Propagation –

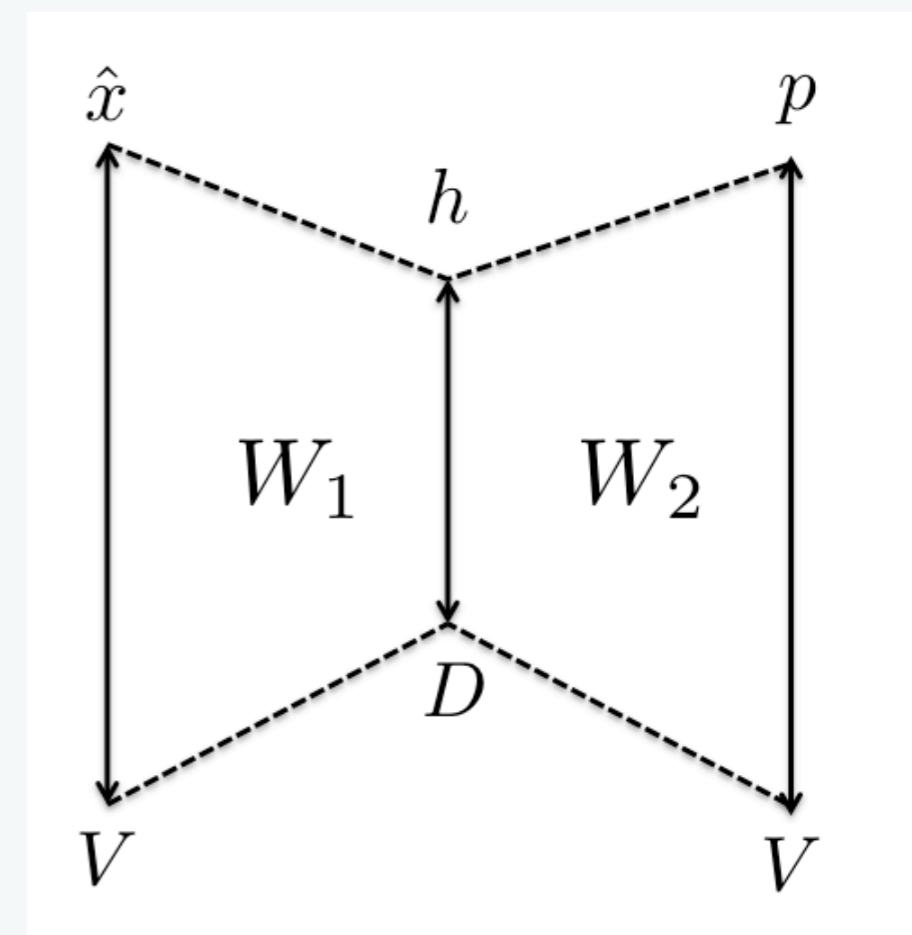
- A one hot vector \hat{x} representing the center word is feeded to the neural network.

- A first linear transformation maps \hat{x} to the D-dimensional vector h as follows:

$$h = W_1^T \hat{x}$$

- A second transformation maps the hidden vector h to the prediction vector p as follows:

$$p = \text{sigmoid}(W_2^T h)$$



The Word2vec approach – The Forward Propagation –

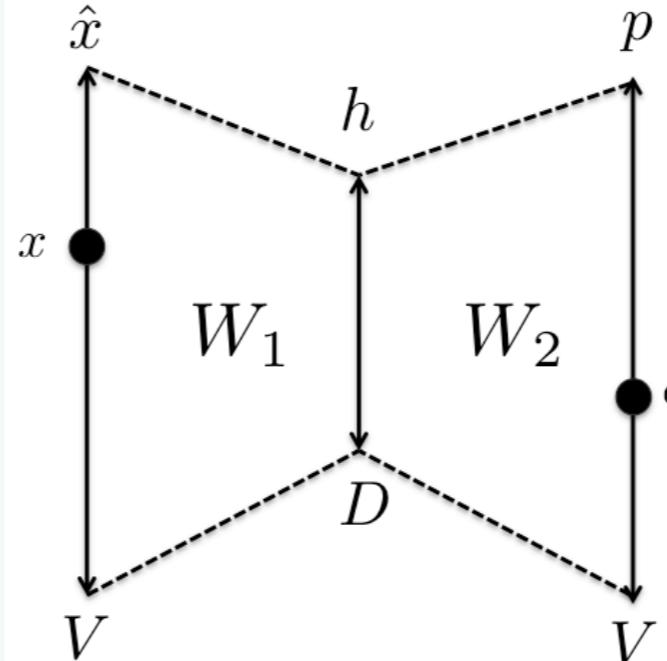
- Let x be the non zero position in the one hot vector \hat{x} and o be one of the V dimensions of the prediction vector p

- We can easily prove that:

$$p_o = \text{sigmoid}(W_1[x]^T W_2[o])$$



- In other words, to predict p_o , we only need the row x of the matrix W_1 and the row o of the matrix W_2

- Let us consider the example of « economy » as a center word in one document, we have the following couples of (center word, context word), for center word in $\mathcal{C} = \{\text{idx}_{\text{for}}, \text{senior}_{\text{Bank}}, \text{idx}_{\text{says}}, \text{idx}_{\text{bosses}}, \text{idx}_{\text{hurting}}\}$



$$W_1 = \begin{bmatrix} W_1[1] \\ W_1[2] \\ \vdots \\ W_1[V] \end{bmatrix}$$

$$W_2 = \begin{bmatrix} | & | & & | \\ W_2[1] & W_2[2] & \cdots & W_2[V] \\ | & | & & | \end{bmatrix}$$



- The loss associated with these 12 samples is:

$$J = - \sum_{c \in \mathcal{C}} [\log(\sigma(W_1[\text{idx}_{\text{economy}}]^T W_2[c])) + \log(1 - \sigma(W_1[\text{idx}_{\text{music}}]^T W_2[c]))]$$

The Word2vec approach – The Learning Process –

Pseudo code:

- Initialize W_1 and W_2 randomly.
- Initialize an empty list of losses.
- For each epoch:
 - Shuffle the sequences.
 - For each sequence in sequences:
 - For each position in the sequence
 - Get the true center word (corresponding to the position).
 - Get the context of the true center word.
 - Get the fake center word.
 - Do one step of SGD for the true center word to update W_1 and W_2
 - Do one step of SGD for the fake center word to update W_1 and W_2
 - Keep track of the loss function by appending the list of losses

Programming Session

